nextwork.org

Connect a Web App to Amazon Aurora

M Mohammed Dawoud Mota

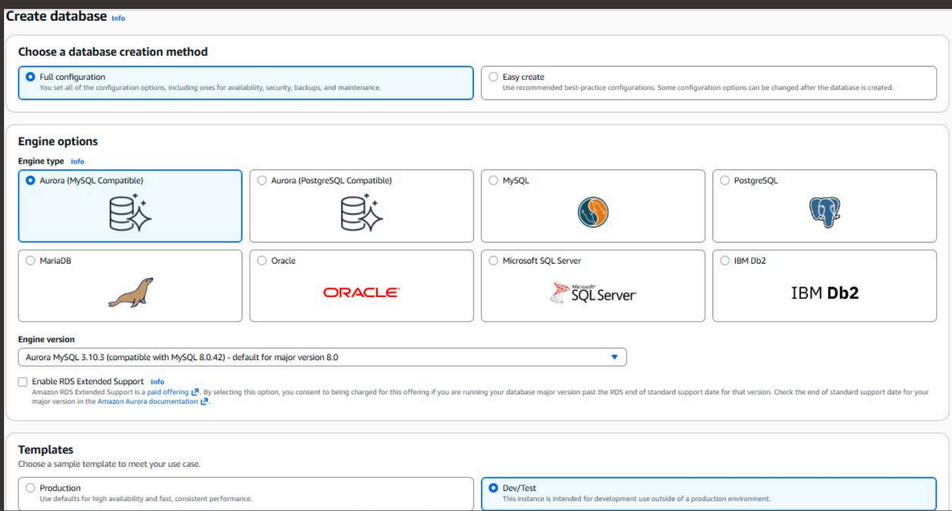

Introducing Today's Project!

What is Amazon Aurora?

Amazon Aurora is a relational database that’s built for high performance and large scale applications. It uses clusters with a writer and reader instance, which keeps your data available and gives you automatic failover and backups. It’s designed for big workloads that need speed, reliability and strong uptime.

How I used Amazon Aurora in this project

I used Amazon Aurora by creating a brand-new Aurora MySQL database cluster and configuring it from scratch. I set the engine type to Aurora MySQL, chose the Dev Test template, created the cluster identifier and admin credentials, and launched the database with a writer and reader instance. I also connected the Aurora database to the EC2 instance I created so it could act as the web app server.

One thing I didn't expect in this project was...

One thing I didn’t expect in this project was seeing how Aurora automatically creates a whole database cluster with a writer and reader instance. I didn’t realize it would show multiple databases in the list right away, and it was interesting to learn that Aurora uses clusters to keep everything highly available and ready for failover.

This project took me...

This project took me around an hour to complete.

M

Mohammed Dawoud M...
NextWork Student

In the first part of my project...

Creating an Aurora Cluster

A relational database stores data in tables with rows and columns, allowing relationships between data points. It uses SQL to easily manage, query, and organize information efficiently.

Aurora is a good choice when you need a relational database to support something large-scale with peak performance and uptime.

Mohammed Dawoud M...
NextWork Student
nextwork.org

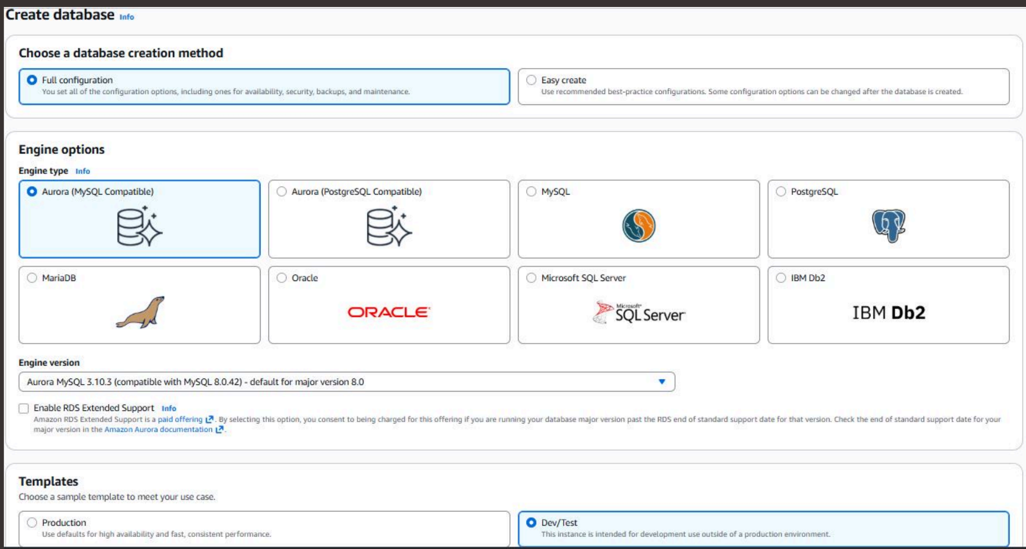

Halfway through I stopped!

I stopped creating my Aurora database because I needed to create an EC2 instance so that I could connect the web app to the database.

Features of my EC2 instance

I created a new key pair for my EC2 instance because this is what will allow me to access my instance and perform any of the needed modifications to the instance.

When I created my EC2 instance, I took particular note of the IPv4 address and the Key pair name. These are important because IPv4 tells you where the server is located, and the key pair is what allows you into the server. Without both, you can not get into the server.

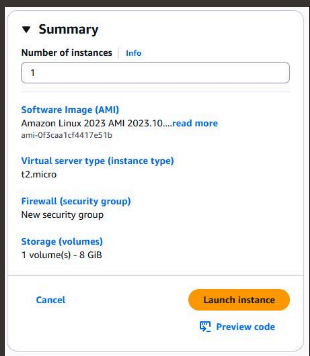

M

Mohammed Dawoud M...
NextWork Student

Then I could finish setting up my database

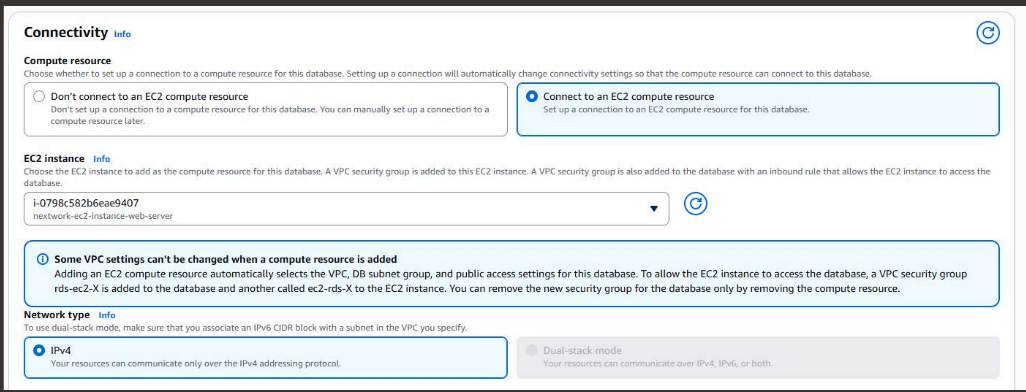

Aurora Database uses clusters because a database cluster in Aurora is a group of database copies that work together, so your data is always available. Each cluster consists of a primary instance (where all write operations occur) and multiple read replicas as back-ups. If your database's primary instance fails, one of the replicas can be promoted to primary automatically.